

# 人工智能基础

## 作业五

### Part1 LSTM RNN GRU 对比试验

#### 1. Recurrent Layers 介绍:

PyTorch提供的常用的循环神经网络层包括以下几种:

- RNN
- GRU
- LSTM

下面我会简单介绍RNN的输入输出的格式、以及初始化参数的含义。

RNN是最基础的循环神经网络, 结构简单。它通过隐藏状态在时间步之间传递信息。其缺点是存在梯度消失/爆炸的问题, 难以学习长序列中的远距离信息关联。

**输入:** 输入通常是序列数据。

- 不使用批量输入 (unbatched) 时, 形状为  $(L, H_{in})$ , 其中  $L$ : 序列长度,  $H_{in}$ : 输入特征数量
- 使用批量输入时, 当 `batch_first=False` 时, 形状为  $(L, N, H_{in})$ ,  $N$ : 批量大小。
- `batch_first=True` 时, 形状为  $(N, L, H_{in})$ 。
- 输入也可以是打包的变长序列 (packed variable length sequence), 可使用 `torch.nn.utils.rnn.pack_padded_sequence()`

**输出:** 所有时间步的隐藏状态。

- 非批量输入时, 形状为  $(L, D * H_{out})$ 。
- 当 `batch_first=False` 时, 形状为  $(L, N, D * H_{out})$ 。
- 当 `batch_first=True` 时, 形状为  $(N, L, D * H_{out})$ 。
- 若输入是 `torch.nn.utils.rnn.PackedSequence`, 输出也会是打包序列。当 `bidirectional=True` 时, `output` 会包含每个时间步的前向和反向隐藏状态的拼接。

**初始化参数含义:**

- `input_size`: 输入序列中每个时间步的特征数量。
- `hidden_size`: 隐藏状态  $h_t$  的特征数量, 它决定了 LSTM 单元在每个时间步处理信息的能力。
- `num_layers`: 循环层的数量, 默认值为 1。例如, `num_layers = 2` 表示将两个 LSTM 层堆叠在一起, 第二层接收第一层的输出并计算最终结果。
- `bias`: 一个布尔值, 若为 `True`, 则 LSTM 层使用偏置权重 `b_ih` 和 `b_hh`; 若为 `False`, 则不使用, 默认值为 `True`。
- `batch_first`: 布尔值, 若为 `True`, 输入和输出张量的维度顺序为  $(batch, seq, feature)$ ; 若为 `False`, 则为  $(seq, batch, feature)$ , 默认值为 `False`。
- `dropout`: 一个浮点数, 默认值为 0。若不为 0, 则在除最后一层外的每个 LSTM 层的输出上引入 Dropout 层, Dropout 概率等于该值。
- `bidirectional`: 布尔值, 若为 `True`, 则使用双向 LSTM, 它可以同时处理序列的正向和反向信息, 默认值为 `False`。

- **proj\_size**: 若大于 0, 则使用带投影的 LSTM。此时, 隐藏状态 `h_t` 的维度会从 `hidden_size` 变为 `proj_size`, 并且每个层的输出隐藏状态会乘以一个可学习的投影矩阵, 默认值为 0。

下面简要介绍 LSTM 和 GRU, 其输入输出和初始化参数的设置可以查阅文档。

- **LSTM**: 一种特殊的 RNN, 通过门控机制解决了传统 RNN 的梯度消失问题, 能有效捕捉长距离依赖。
- **GRU**: LSTM 的简化版, 合并了一些门控机制, 提高了计算效率。

## 2. 运行实验结果

### LSTM

```
vocab_size: 20001
ImdbNet(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM(64, 64)
  (linear1): Linear(in_features=64, out_features=64, bias=True)
  (act1): ReLU()
  (linear2): Linear(in_features=64, out_features=2, bias=True)
)
Train Epoch: 1 Loss: 0.583308    Acc: 0.674171
Test set: Average loss: 0.4625, Accuracy: 0.7801
Train Epoch: 2 Loss: 0.388254    Acc: 0.824231
Test set: Average loss: 0.3808, Accuracy: 0.8321
Train Epoch: 3 Loss: 0.298237    Acc: 0.875649
Test set: Average loss: 0.3619, Accuracy: 0.8475
Train Epoch: 4 Loss: 0.236970    Acc: 0.906799
Test set: Average loss: 0.3437, Accuracy: 0.8580
Train Epoch: 5 Loss: 0.184549    Acc: 0.932907
Test set: Average loss: 0.3604, Accuracy: 0.8588
```

### RNN

```
vocab_size: 20001
ImdbNet(
  (embedding): Embedding(20001, 64)
  (rnn): RNN(64, 64)
  (linear1): Linear(in_features=64, out_features=64, bias=True)
  (act1): ReLU()
  (linear2): Linear(in_features=64, out_features=2, bias=True)
)
Train Epoch: 1 Loss: 0.608334    Acc: 0.651158
Test set: Average loss: 0.5046, Accuracy: 0.7545
Train Epoch: 2 Loss: 0.417458    Acc: 0.814846
Test set: Average loss: 0.4019, Accuracy: 0.8159
Train Epoch: 3 Loss: 0.322093    Acc: 0.867961
Test set: Average loss: 0.3807, Accuracy: 0.8331
Train Epoch: 4 Loss: 0.259338    Acc: 0.898762
Test set: Average loss: 0.3517, Accuracy: 0.8507
Train Epoch: 5 Loss: 0.208925    Acc: 0.922125
Test set: Average loss: 0.3815, Accuracy: 0.8499
```

### GRU

```
vocab_size: 20001
ImdbNet(
```

```

(embedding): Embedding(20001, 64)
(gru): GRU(64, 64)
(linear1): Linear(in_features=64, out_features=64, bias=True)
(act1): ReLU()
(linear2): Linear(in_features=64, out_features=2, bias=True)
)
Train Epoch: 1 Loss: 0.571737    Acc: 0.687899
Test set: Average loss: 0.4627, Accuracy: 0.7811
Train Epoch: 2 Loss: 0.383377    Acc: 0.827676
Test set: Average loss: 0.3747, Accuracy: 0.8325
Train Epoch: 3 Loss: 0.303219    Acc: 0.869758
Test set: Average loss: 0.3621, Accuracy: 0.8443
Train Epoch: 4 Loss: 0.237048    Acc: 0.905950
Test set: Average loss: 0.3356, Accuracy: 0.8566
Train Epoch: 5 Loss: 0.184019    Acc: 0.932109
Test set: Average loss: 0.3499, Accuracy: 0.8570

```

总结: Acc均在85%左右, 其中RNN相对偏低(与理论相符), 总体无明显差异。

## Part2 手写LSTM实验

### 1. 超过80%实验结果: (代码后附)

```

隐藏层维度: 128; 损失函数: CrossEntropyLoss; Epoch: 10; batch_size: 64
Test set: Average loss: 0.6099, Accuracy: 0.8317

```

### 2. 调整过程

#### 1. 原始结构: 隐藏层维度: 64; 损失函数: CrossEntropyLoss; Epoch: 5; batch\_size: 64

```

Net(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM()
  (fc1): Linear(in_features=64, out_features=64, bias=True)
  (fc2): Linear(in_features=64, out_features=2, bias=True)
)
100%|████████████████████████████████████████| 313/313 [01:05<00:00, 4.76it/s]
Train Epoch: 1 Loss: 0.693271    Acc: 0.505741
Test set: Average loss: 0.6912, Accuracy: 0.5196
100%|████████████████████████████████████████| 313/313 [01:06<00:00, 4.68it/s]
Train Epoch: 2 Loss: 0.681398    Acc: 0.551967
Test set: Average loss: 0.6684, Accuracy: 0.5833
100%|████████████████████████████████████████| 313/313 [01:07<00:00, 4.61it/s]
Train Epoch: 3 Loss: 0.632539    Acc: 0.645118
Test set: Average loss: 0.6545, Accuracy: 0.6185
100%|████████████████████████████████████████| 313/313 [01:08<00:00, 4.54it/s]
Train Epoch: 4 Loss: 0.572839    Acc: 0.713409
Test set: Average loss: 0.5932, Accuracy: 0.7065
100%|████████████████████████████████████████| 313/313 [01:06<00:00, 4.72it/s]
Train Epoch: 5 Loss: 0.483343    Acc: 0.783846
Test set: Average loss: 0.6495, Accuracy: 0.6406

```

发现: 准确率不高。

#### 2. 调整网络结构

- 增加一层全连接层: 隐藏层维度: 64; 损失函数: CrossEntropyLoss; Epoch: 5; batch\_size: 64

```

(embedding): Embedding(20001, 64)
(lstm): LSTM()
(fc1): Linear(in_features=64, out_features=64, bias=True)
(fc2): Linear(in_features=64, out_features=64, bias=True)
(fc3): Linear(in_features=64, out_features=2, bias=True)
)
100%|████████████████████████████████████████| 313/313 [01:05<00:00, 4.81it/s]
Train Epoch: 1 Loss: 0.693682 Acc: 0.501398
Test set: Average loss: 0.6929, Accuracy: 0.5077
100%|████████████████████████████████████████| 313/313 [01:04<00:00, 4.86it/s]
Train Epoch: 2 Loss: 0.685300 Acc: 0.537590
Test set: Average loss: 0.7010, Accuracy: 0.5053
100%|████████████████████████████████████████| 313/313 [01:08<00:00, 4.58it/s]
Train Epoch: 3 Loss: 0.672419 Acc: 0.562899
Test set: Average loss: 0.6691, Accuracy: 0.5821
100%|████████████████████████████████████████| 313/313 [01:10<00:00, 4.45it/s]
Train Epoch: 4 Loss: 0.660106 Acc: 0.579473
Test set: Average loss: 0.6862, Accuracy: 0.5061
100%|████████████████████████████████████████| 313/313 [01:05<00:00, 4.75it/s]
Train Epoch: 5 Loss: 0.682045 Acc: 0.538488

Test set: Average loss: 0.6881, Accuracy: 0.5425

```

发现：Acc很低，离80%的标准距离很远。

- **隐藏层维度：128**；损失函数：CrossEntropyLoss；Epoch: 10；batch\_size: 64

```

Net(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM()
  (fc1): Linear(in_features=128, out_features=128, bias=True)
  (fc2): Linear(in_features=128, out_features=2, bias=True)
)
100%|████████████████████████████████████████| 313/313 [01:06<00:00, 4.72it/s]
Train Epoch: 1 Loss: 0.692490 Acc: 0.513279
Test set: Average loss: 0.6869, Accuracy: 0.5471
100%|████████████████████████████████████████| 313/313 [01:02<00:00, 4.97it/s]
Train Epoch: 2 Loss: 0.679165 Acc: 0.552915
Test set: Average loss: 0.6734, Accuracy: 0.5987
100%|████████████████████████████████████████| 313/313 [01:04<00:00, 4.86it/s]
Train Epoch: 3 Loss: 0.656441 Acc: 0.590755
Test set: Average loss: 0.6955, Accuracy: 0.5101
100%|████████████████████████████████████████| 313/313 [01:06<00:00, 4.73it/s]
Train Epoch: 4 Loss: 0.645554 Acc: 0.620457
Test set: Average loss: 0.5636, Accuracy: 0.7296
100%|████████████████████████████████████████| 313/313 [01:06<00:00, 4.71it/s]
Train Epoch: 5 Loss: 0.426162 Acc: 0.812899
Test set: Average loss: 0.4266, Accuracy: 0.8093
100%|████████████████████████████████████████| 313/313 [01:03<00:00, 4.95it/s]
Train Epoch: 6 Loss: 0.301800 Acc: 0.877895
Test set: Average loss: 0.3987, Accuracy: 0.8285

```

```

100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████|

313/313 [01:05<00:00, 4.80it/s]
Train Epoch: 7 Loss: 0.225721 Acc: 0.916334
Test set: Average loss: 0.4058, Accuracy: 0.8426
100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████|

313/313 [01:03<00:00, 4.91it/s]
Train Epoch: 8 Loss: 0.166375 Acc: 0.942792
Test set: Average loss: 0.4409, Accuracy: 0.8384
100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████|

313/313 [01:04<00:00, 4.86it/s]
Train Epoch: 9 Loss: 0.110612 Acc: 0.965605
Test set: Average loss: 0.4804, Accuracy: 0.8398
100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████|

313/313 [01:05<00:00, 4.81it/s]
Train Epoch: 10 Loss: 0.068540 Acc: 0.980831
Test set: Average loss: 0.6099, Accuracy: 0.8317

```

发现：增加隐藏层维度可以明显提升测试集Acc，在5个epoch结束后已经达到80%且后续训练可以继续提升准确率。

### 3. 调整损失函数

- PyTorch官方文档给出的几种损失函数中CrossEntropyLoss是最适合用于此问题的，其他不合适的损失函数可能会造成一些严重后果（如CTCLoss多用于处理序列长度可变的任务，如语音识别。Acc达到了0%，因此我直接中止了程序运行）

隐藏层维度：64；损失函数：CTCLoss；Epoch：5（提前终止）；batch\_size：64

```

Net(
  (embedding): Embedding(20000, 64)
  (lstm): LSTM()
  (fc): Linear(in_features=64, out_features=3, bias=True)
)
100%|████████████████████████████████████████████████████████████████████████████████| 391/391 [01:16<00:00, 5.11it/s]
Train Epoch: 1 Loss: 10.216603 Acc: 0.001279
Test set: Average loss: 1.4054, Accuracy: 0.0000
100%|████████████████████████████████████████████████████████████████████████████████| 391/391 [01:22<00:00, 4.73it/s]
Train Epoch: 2 Loss: 1.353754 Acc: 0.000000
Traceback (most recent call last):

```

发现：损失函数对训练影响非常大，要根据目的选择合适的损失函数。

- 通过网络搜索到了Label Smoothing Cross Entropy 损失函数。它有助于防止过拟合并减轻错误标签的影响。

隐藏层维度：64；损失函数：**Label Smoothing Cross Entropy Loss**；Epoch：5；batch\_size：64

使用设备：cuda:0

```

Net(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM()
  (fc1): Linear(in_features=64, out_features=64, bias=True)
  (fc2): Linear(in_features=64, out_features=2, bias=True)
)

```

```

)
Train Epoch: 1 Loss: 0.691385 Acc: 0.522214: 100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
| 313/313 [01:08<00:00, 4.55it/s]
Train Epoch: 1 Loss: 0.691385 Acc: 0.522214
Test set: Average loss: 0.6861, Accuracy: 0.5500
Train Epoch: 2 Loss: 0.684264 Acc: 0.551018: 100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
| 313/313 [01:05<00:00, 4.78it/s]
Train Epoch: 2 Loss: 0.684264 Acc: 0.551018
Test set: Average loss: 0.6904, Accuracy: 0.5320
Train Epoch: 3 Loss: 0.650201 Acc: 0.639227: 100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
| 313/313 [01:09<00:00, 4.50it/s]
Train Epoch: 3 Loss: 0.650201 Acc: 0.639227
Test set: Average loss: 0.5966, Accuracy: 0.7217
Train Epoch: 4 Loss: 0.533848 Acc: 0.771915: 100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
| 313/313 [01:07<00:00, 4.64it/s]
Train Epoch: 4 Loss: 0.533848 Acc: 0.771915
Test set: Average loss: 0.5130, Accuracy: 0.7836
Train Epoch: 5 Loss: 0.443213 Acc: 0.843251: 100%|
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
████████████████████████████████████████████████████████████████████████████████
| 313/313 [01:07<00:00, 4.63it/s]
Train Epoch: 5 Loss: 0.443213 Acc: 0.843251
Test set: Average loss: 0.4712, Accuracy: 0.8192

```

发现：表现明显优于只使用CrossEntropyLoss，在隐藏层维度64的前提下仍然达到了80+的准确率。过拟合现象被抑制。

4. 调整训练流程--如超参数，epoch，batchsize等：

- 隐藏层维度：64；损失函数：CrossEntropyLoss；**Epoch: 10**；batch\_size: 64

```

Net(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM()
  (fc1): Linear(in_features=64, out_features=64, bias=True)
  (fc2): Linear(in_features=64, out_features=2, bias=True)
)
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:06<00:00, 4.70it/s]
Train Epoch: 1 Loss: 0.693224 Acc: 0.507937
Test set: Average loss: 0.6923, Accuracy: 0.5184
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:00<00:00, 5.20it/s]
Train Epoch: 2 Loss: 0.689369 Acc: 0.528754
Test set: Average loss: 0.6905, Accuracy: 0.5182
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:01<00:00, 5.13it/s]
Train Epoch: 3 Loss: 0.677964 Acc: 0.558956
Test set: Average loss: 0.6757, Accuracy: 0.5698
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:03<00:00, 4.96it/s]
Train Epoch: 4 Loss: 0.652045 Acc: 0.594399
Test set: Average loss: 0.6662, Accuracy: 0.6001
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:03<00:00, 4.95it/s]
Train Epoch: 5 Loss: 0.635427 Acc: 0.606679
Test set: Average loss: 0.6698, Accuracy: 0.5932
100%|████████████████████████████████████████████████████████████████████████████████| 313/313 [01:02<00:00, 5.01it/s]
Train Epoch: 6 Loss: 0.562416 Acc: 0.697434

```

```

Test set: Average loss: 0.5929, Accuracy: 0.7004
100%|████████████████████████████████████████| 313/313 [01:02<00:00, 5.00it/s]
Train Epoch: 7 Loss: 0.557378 Acc: 0.705970
Test set: Average loss: 0.5609, Accuracy: 0.7373
100%|████████████████████████████████████████| 313/313 [01:02<00:00, 5.01it/s]
Train Epoch: 8 Loss: 0.427962 Acc: 0.816693
Test set: Average loss: 0.5506, Accuracy: 0.7579
100%|████████████████████████████████████████| 313/313 [01:02<00:00, 5.01it/s]
Train Epoch: 9 Loss: 0.348632 Acc: 0.858826
Test set: Average loss: 0.5134, Accuracy: 0.7987
100%|████████████████████████████████████████| 313/313 [01:03<00:00, 4.93it/s]
Train Epoch: 10 Loss: 0.287051 Acc: 0.891074
Test set: Average loss: 0.5157, Accuracy: 0.7975

```

发现：增加循环次数可以提高准确率（在未收敛时），同时也会增加时间成本。但要注意，过度增加 epoch 可能会导致过拟合，即模型在训练集上表现很好，但在测试集或新数据上表现不佳。

- 隐藏层维度：64；损失函数：CrossEntropyLoss；Epoch：10；**batch\_size：32**

```

Net(
  (embedding): Embedding(20001, 64)
  (lstm): LSTM()
  (fc1): Linear(in_features=64, out_features=64, bias=True)
  (fc2): Linear(in_features=64, out_features=2, bias=True)
)
100%|████████████████████████████████████████| 625/625 [02:06<00:00, 4.95it/s]
Train Epoch: 1 Loss: 0.693041 Acc: 0.512000
Test set: Average loss: 0.6903, Accuracy: 0.5303
100%|████████████████████████████████████████| 625/625 [02:08<00:00, 4.88it/s]
Train Epoch: 2 Loss: 0.676670 Acc: 0.562550
Test set: Average loss: 0.6718, Accuracy: 0.5961
100%|████████████████████████████████████████| 625/625 [02:07<00:00, 4.90it/s]
Train Epoch: 3 Loss: 0.606563 Acc: 0.680600
Test set: Average loss: 0.6091, Accuracy: 0.6865
100%|████████████████████████████████████████| 625/625 [02:03<00:00, 5.06it/s]
Train Epoch: 4 Loss: 0.489389 Acc: 0.781200
Test set: Average loss: 0.6028, Accuracy: 0.7402
100%|████████████████████████████████████████| 625/625 [02:05<00:00, 4.98it/s]
Train Epoch: 5 Loss: 0.405098 Acc: 0.832950
Test set: Average loss: 0.5037, Accuracy: 0.7739
100%|████████████████████████████████████████| 625/625 [02:07<00:00, 4.90it/s]
Train Epoch: 6 Loss: 0.348666 Acc: 0.860550
Test set: Average loss: 0.4871, Accuracy: 0.8035
100%|████████████████████████████████████████| 625/625 [02:11<00:00, 4.75it/s]
Train Epoch: 7 Loss: 0.281200 Acc: 0.892750
Test set: Average loss: 0.5112, Accuracy: 0.7846
100%|████████████████████████████████████████| 625/625 [02:06<00:00, 4.93it/s]
Train Epoch: 8 Loss: 0.259462 Acc: 0.905250
Test set: Average loss: 0.7532, Accuracy: 0.5018
100%|████████████████████████████████████████| 625/625 [02:07<00:00, 4.89it/s]
Train Epoch: 9 Loss: 0.332733 Acc: 0.852900
Test set: Average loss: 0.5388, Accuracy: 0.7976
100%|████████████████████████████████████████| 625/625 [02:10<00:00, 4.79it/s]
Train Epoch: 10 Loss: 0.214364 Acc: 0.919700
Test set: Average loss: 0.5690, Accuracy: 0.8145

```

发现：Acc比batch\_size=64时有小幅的提高，可能的原因是：较小的 batch size 意味着模型在训练过程中看到的数据更多样化，每次更新参数时使用的样本不同，这有助于模型学习到更通用的特征，从而提高泛化能力。但模型训练稳定性可能会下降。

## 特别鸣谢

图书馆西区一层的高性能工作区，帮我节约了大量训练时间。

## code

全部代码已附在压缩包中。

手写实现lstm代码并达到80%+Acc代码见下：

```
# 用Pytorch手写一个LSTM网络，在IMDB数据集上进行训练
# 跑5个epoch，hidden size保持64不变
import os
import numpy as np
import torch
import torch.nn as nn

from torch.utils.data import Dataset, DataLoader
from utils import load_imdb_dataset, Accuracy
import sys

from utils import load_imdb_dataset, Accuracy
from tqdm import tqdm
import math

# 超参数
n_epoch = 10
batch_size = 64
print_freq = 2

# print("MLU is not available, use GPU/CPU instead.")
if torch.cuda.is_available():
    device = torch.device('cuda:0')
else:
    device = torch.device('cpu')

X_train, y_train, X_test, y_test = load_imdb_dataset('data', nb_words=20000,
test_split=0.2)

seq_len = 200
vocab_size = len(X_train) + 1

class ImdbDataset(Dataset):

    def __init__(self, x, y):
        self.x = x
```



```

        self.y = y

    def __getitem__(self, index):

        data = self.x[index]
        data = np.concatenate([data[:seq_Len], [0] * (seq_Len -
len(data))]).astype('int32') # set
        label = self.y[index]
        return data, label

    def __len__(self):

        return len(self.y)

# 你需要实现的手写LSTM内容，包括LSTM类所属的__init__函数和forward函数
class LSTM(nn.Module):
    """
    手写lstm，可以用全连接层nn.Linear，不能直接用nn.LSTM
    """
    def __init__(self, input_size, hidden_size):
        super(LSTM, self).__init__()
        self.hidden_size = hidden_size
        self.imput_size = input_size

        # 初始化所有门控的参数：权重w和偏置b

        # 输入门
        self.w_ii = nn.Parameter(torch.Tensor(hidden_size, input_size))
        self.w_hi = nn.Parameter(torch.Tensor(hidden_size, hidden_size))
        self.b_ii = nn.Parameter(torch.Tensor(hidden_size))

        # 遗忘门
        self.w_if = nn.Parameter(torch.Tensor(hidden_size, input_size))
        self.w_hf = nn.Parameter(torch.Tensor(hidden_size, hidden_size))
        self.b_if = nn.Parameter(torch.Tensor(hidden_size))

        # 细胞状态
        self.w_ic = nn.Parameter(torch.Tensor(hidden_size, input_size))
        self.w_hc = nn.Parameter(torch.Tensor(hidden_size, hidden_size))
        self.b_ic = nn.Parameter(torch.Tensor(hidden_size))

        # 输出门
        self.w_io = nn.Parameter(torch.Tensor(hidden_size, input_size))
        self.w_ho = nn.Parameter(torch.Tensor(hidden_size, hidden_size))
        self.b_io = nn.Parameter(torch.Tensor(hidden_size))

        # 初始化参数
        self.reset_parameters()

    def reset_parameters(self):
        # 使用均匀分布初始化权重。常用策略
        stdv = 1.0 / math.sqrt(self.hidden_size)
        for weight in self.parameters():
            nn.init.uniform_(weight, -stdv, stdv)

```

```

def forward(self, x, init_states=None):
    """
    x: 输入序列, shape: (batch_size, seq_len, input_size)
    返回: output, (h_n, c_n)
    """
    batch_size, seq_len, _ = x.size()

    # 初始化隐藏状态和细胞状态
    if init_states is None:
        h_t = torch.zeros(batch_size, self.hidden_size).to(x.device)
        c_t = torch.zeros(batch_size, self.hidden_size).to(x.device)
    else:
        h_t, c_t = init_states

    # 存储所有时间步的输出
    outputs = []

    for t in range(seq_len):
        # 获取当前时间步的输入
        x_t = x[:, t, :] # (batch_size, input_size)

        # 输入门计算
        i_t = torch.sigmoid(
            torch.matmul(x_t, self.w_ii.t()) + torch.matmul(h_t, self.w_hi.t()) +
self.b_ii
        )

        # 遗忘门计算
        f_t = torch.sigmoid(
            torch.matmul(x_t, self.w_if.t()) + torch.matmul(h_t, self.w_hf.t()) +
self.b_if
        )

        # 细胞状态计算
        c_tilde = torch.tanh(
            torch.matmul(x_t, self.w_ic.t()) + torch.matmul(h_t, self.w_hc.t()) +
self.b_ic
        )

        # 更新细胞状态
        c_t = f_t * c_t + i_t * c_tilde

        # 输出门计算
        o_t = torch.sigmoid(
            torch.matmul(x_t, self.w_io.t()) + torch.matmul(h_t, self.w_ho.t()) +
self.b_io
        )

        # 更新隐藏状态
        h_t = o_t * torch.tanh(c_t)

        # 保存当前时间步的输出
        outputs.append(h_t.unsqueeze(1))

    # 拼接所有时间步的输出
    outputs = torch.cat(outputs, dim=1) # (batch_size, seq_len, hidden_size)

```

```
    return outputs, (h_t, c_t)
```

```
class Net(nn.Module):
```

```
    '''
```

```
    一层LSTM的文本分类模型
```

```
    '''
```

```
    def __init__(self, embedding_size=64, hidden_size=128, num_classes=2):
```

```
        super(Net, self).__init__()
```

```
        self.embedding = nn.Embedding(vocab_size, embedding_size)
```

```
        self.lstm = LSTM(input_size=embedding_size, hidden_size=hidden_size)
```

```
        self.fc1 = nn.Linear(hidden_size, hidden_size)
```

```
        self.fc2 = nn.Linear(hidden_size, num_classes)
```

```
    def forward(self, x):
```

```
        '''
```

```
        x: 输入, shape: (batch_size, seq_len)
```

```
        '''
```

```
        # 词嵌入 (batch_size, seq_len) -> (batch_size, seq_len, embedding_size)
```

```
        x = self.embedding(x)
```

```
        # LSTM层 (batch_size, seq_len, embedding_size) -> (batch_size, seq_len, hidden_size)
```

```
        lstm_out, _ = self.lstm(x)
```

```
        # 取最后一个时间步的输出 (batch_size, hidden_size)
```

```
        last_out = lstm_out[:, -1, :]
```

```
        # 全连接层
```

```
        x = torch.relu(self.fc1(last_out))
```

```
        x = self.fc2(x)
```

```
    return x
```

```
train_dataset = ImdbDataset(X=X_train, y=y_train)
```

```
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
```

```
test_dataset = ImdbDataset(X=X_test, y=y_test)
```

```
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)
```

```
net = Net()
```

```
metric = Accuracy()
```

```
print(net)
```

```
def train(model, device, train_loader, optimizer, epoch):
```

```
    model = model.to(device)
```

```
    model.train()
```

```
    loss_func = torch.nn.CrossEntropyLoss(reduction="mean")
```

```
    train_acc = 0
```

```
    train_loss = 0
```

```
    n_iter = 0
```

```
    for batch_idx, (data, target) in tqdm(enumerate(train_loader),
```

```
total=len(train_loader)):
```

```
        target = target.long()
```

```
        data, target = data.to(device), target.to(device)
```

```

optimizer.zero_grad()
output = model(data)
# loss = F.nll_loss(output, target)
loss = loss_func(output, target)
loss.backward()
optimizer.step()
metric.update(output, target)
train_acc += metric.result()
train_loss += loss.item()
metric.reset()
n_iter += 1

print('Train Epoch: {} Loss: {:.6f} \t Acc: {:.6f}'.format(epoch, train_loss / n_iter,
train_acc / n_iter))

def test(model, device, test_loader):
    model = model.to(device)
    model.eval()
    loss_func = torch.nn.CrossEntropyLoss(reduction="mean")
    test_loss = 0
    test_acc = 0
    n_iter = 0
    with torch.no_grad():
        for data, target in test_loader:
            target = target.long()
            data, target = data.to(device), target.to(device)
            output = model(data)
            # test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up
batch loss
            test_loss += loss_func(output, target).item()
            metric.update(output, target)
            test_acc += metric.result()
            metric.reset()
            n_iter += 1
    test_loss /= n_iter
    test_acc /= n_iter
    print('Test set: Average loss: {:.4f}, Accuracy: {:.4f}'.format(test_loss, test_acc))

optimizer = torch.optim.Adam(net.parameters(), lr=1e-3, weight_decay=0.0)
gamma = 0.7
for epoch in range(1, n_epoch + 1):
    train(net, device, train_loader, optimizer, epoch)
    test(net, device, test_loader)

```



